

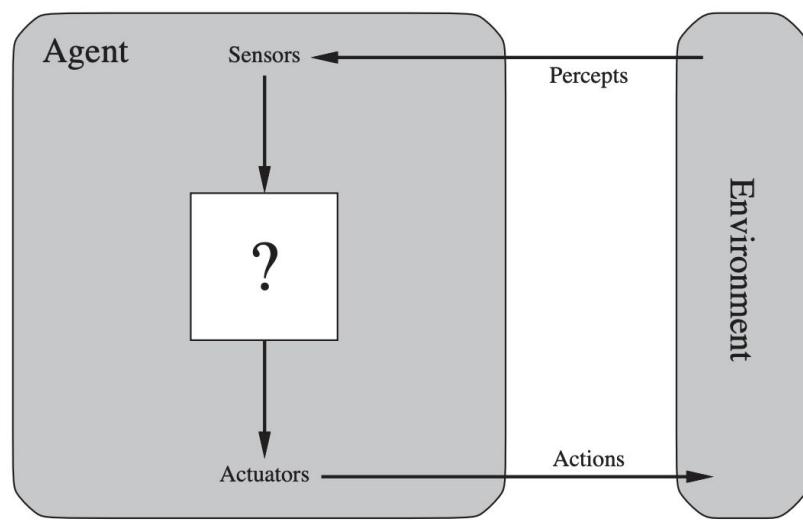
AI Course

Intelligent Agents and Search

Dr. Müsel Taşgın

Agents

Intelligent Agents

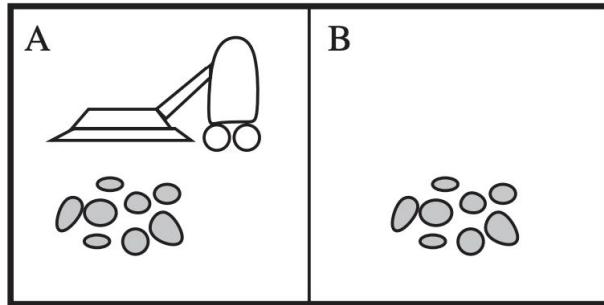


percept: the agent's perceptual inputs at any given instant

percept sequence: the complete history of everything the agent has ever perceived

agent function maps any given percept sequence to an action

Intelligent Agents - *Vacuum cleaner*



Percept sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
:	:
[A, Clean], [A, Clean], [A, Clean]	Right
[A, Clean], [A, Clean], [A, Dirty]	Suck
:	:

A **rational agent** is one that does the right thing—conceptually speaking, every entry in the table for the agent function is filled out correctly. Obviously, doing the right thing is better than doing the wrong thing, but what does it mean to do the right thing?

Types of agents

Simple reflex agent

- No “state” or memory
- Reacts to current input according to its program (rules of the form “if condition then action”)

Model-based agent

- Uses an explicit knowledge base to model the environment
 - How does the environment evolve independently
 - How does the agent affect the environment
- Exhibits “understanding” of its input by relating it to prior knowledge
- Reacts according to rules
 - Conditions may be complex and require inference to evaluate

Types of agents

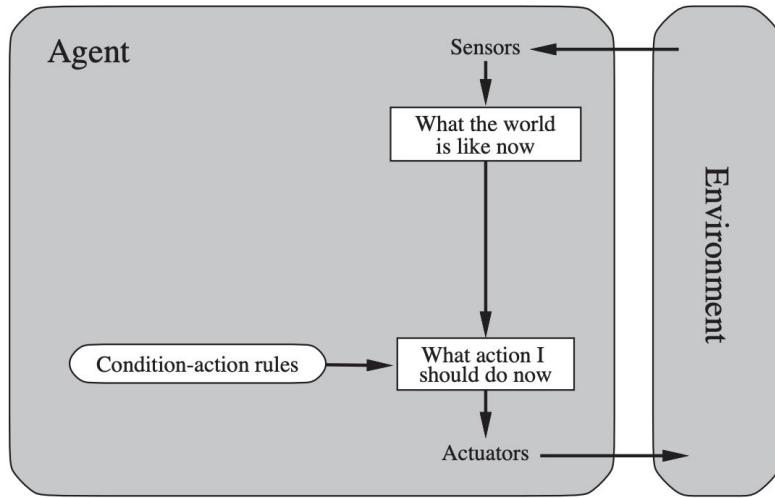
Planning agents (*goal*-based and *utility*-based agents)

- Explicitly represent their own goals and/or preferences (“utilities”) and can reason about them (i.e., planning)
- **Exhibit autonomy:** Actions do not follow directly from rule-based lookup

Learning agents

- **Supervised learning:** Learn from positive and negative examples
- **Reinforcement learning:** Learn from experience to improve its outcomes

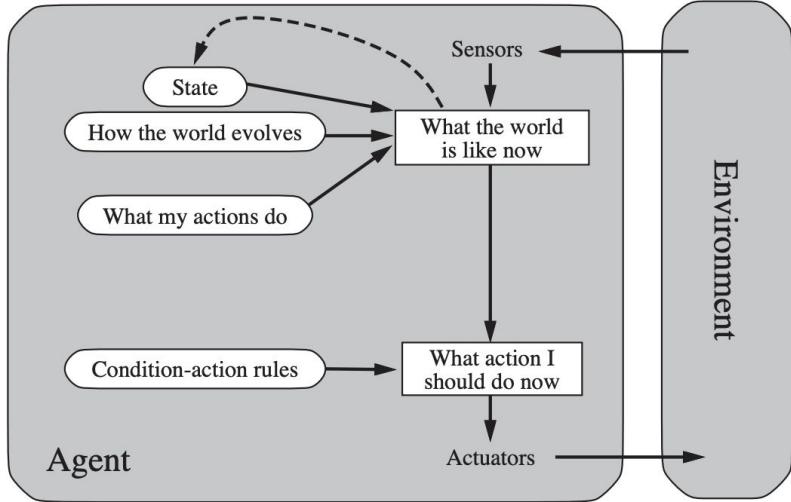
Declarative simple reflex agents



```
function SIMPLE-REFLEX-AGENT(percept) returns an action  
persistent: rules, a set of condition-action rules
```

```
state  $\leftarrow$  INTERPRET-INPUT(percept)  
rule  $\leftarrow$  RULE-MATCH(state, rules)  
action  $\leftarrow$  rule.ACTION  
return action
```

Model-based reflex agents

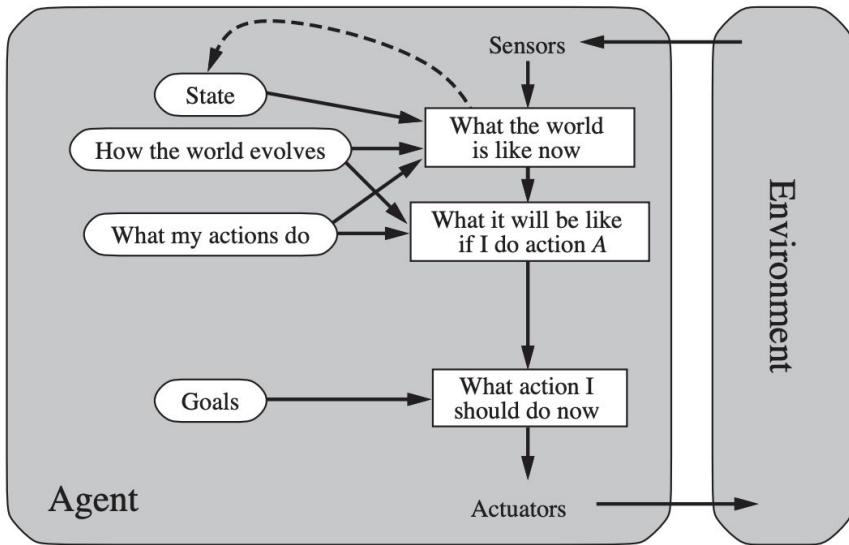


Model: The definition of the environment, e agent's model of how the world works.

```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
  persistent: state, the agent's current conception of the world state
            model, a description of how the next state depends on current state and action
            rules, a set of condition-action rules
            action, the most recent action, initially none

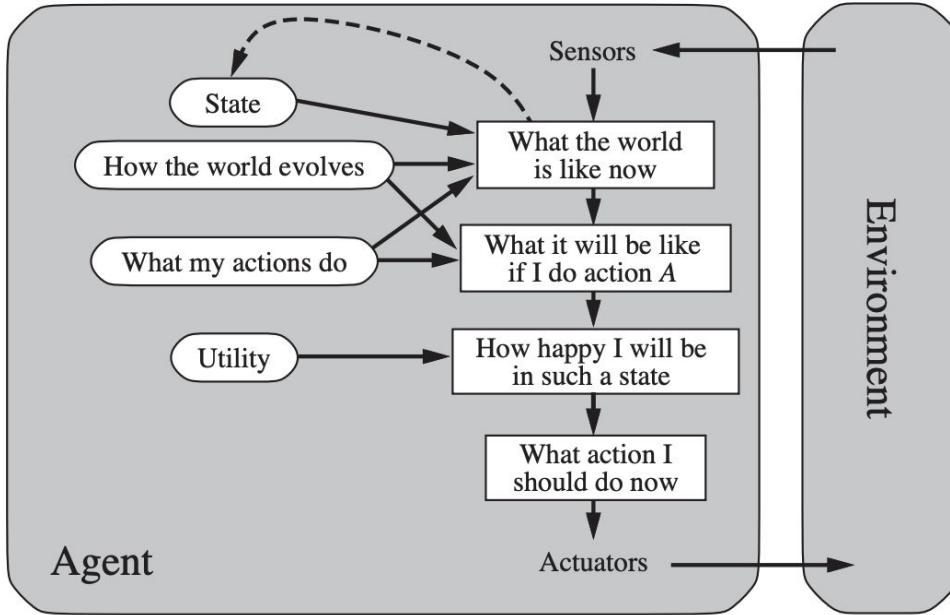
  state  $\leftarrow$  UPDATE-STATE(state, action, percept, model)
  rule  $\leftarrow$  RULE-MATCH(state, rules)
  action  $\leftarrow$  rule.ACTION
  return action
```

Goal-based agents



The agent needs some sort of **goal** information that describes situations that are desirable, for example, being at the passenger's destination.

Utility-based agents



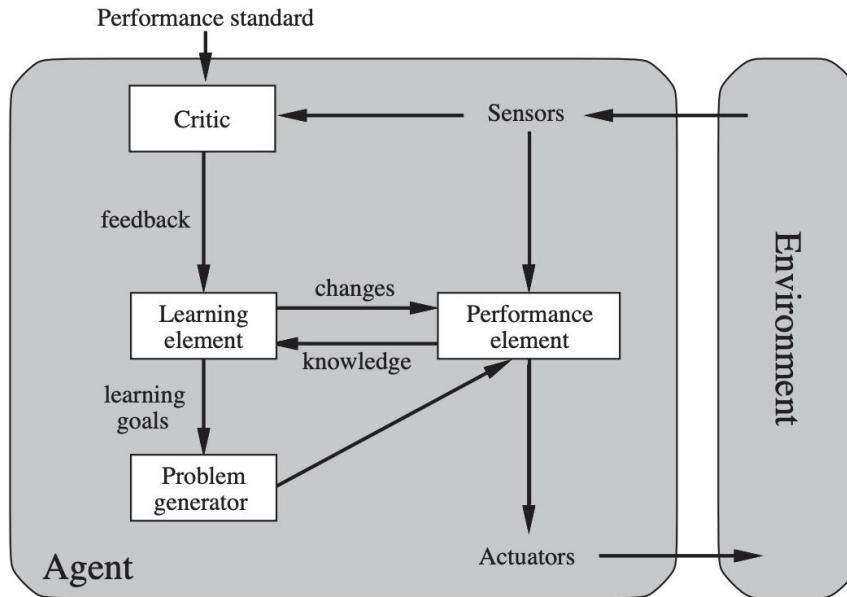
The agent's preferences are captured by a **utility function**, $U(s)$, which assigns a single number to express desirability of a state.

Rational agent tries to select the best action to maximize the utility

Goals alone are not enough to generate high-quality behavior in most environments. For example, many action sequences will get the taxi to its destination (thereby achieving the goal) but some are **quicker**, **safer**, **more reliable**, or **cheaper** than others.

Goals just provide a crude binary distinction between "**happy**" and "**unhappy**" states. A more general performance measure should allow a comparison of different world states according to exactly how happy they would make the agent. Because "happy" does not sound very scientific, economists and computer scientists use the term **utility** instead

Learning agents



A learning agent can be divided into four conceptual components, as shown in figure.

The most important distinction is between the **learning** element, which is responsible for making improvements, and the **performance** element, which is responsible for selecting external actions.

The **performance** element is what we have previously considered to be the entire agent: it takes in percepts and decides on actions.

The **learning** element uses feedback from the critic on how the agent is doing and determines how the **performance** element should be modified to do better in the future.

Drawbacks of simple reflex agents

- **HUGE** rule base, time consuming to build by hand
- What if more than one **condition** is satisfied?
- Inflexible (no **adaptation** or **learning**)

Representing agent knowledge

- **Q:** What formal language(s) can we use to represent
 1. Current facts about the state of the world?
 2. A model of how the world behaves?
 3. A model of the effects of actions that the agent can perform?
 4. The production rules that specify agent behavior?
- **A:** Formal logic
 - Syntax and semantics are well understood
 - Computational tractability known for important subsets (e.g., Horn clause logic)

How do we determine success of the agent?

1. Getting the “*right answer*”?
2. The *Turing Test* (or modified versions)?
3. Having a good outcome? (*using some “utility” function*)
4. Know it when *we see* it?

Analyzing agent performance

- **Rational agent** is one that does the “right” thing
 - Must define a **performance measure**
 - Costs (penalties) and rewards
 - Chooses an action that maximizes expected score
- Rationality depends on
 1. Success criterion defined by performance measure
 2. “Behavior” of the environment (e.g., can a clean square get dirty again?)
 3. Possible actions
 4. Percept sequence
- Autonomy
 - Rational agents require **learning** to compensate for incorrect or incomplete starting knowledge

Rational Agent

Without loss of generality, “goals” specifiable by performance measure defining a numerical value for any environment history

Rational action: whichever action maximizes the expected value of the performance measure given the percept sequence to date

Rational \neq omniscient

Rational \neq clairvoyant

Rational \neq successful

What is rational at any given time depends on four things:

- The **performance measure** that defines the criterion of **success**.
- The agent’s **prior knowledge** of the environment.
- The **actions** that the agent can perform.
- The agent’s **percept sequence** to date.

Rational Agent, Autonomy, Exploration

- **Autonomy:** Learn and act (percepts)
- **Exploration:** Learning new things, trial
- **Exploitation:** Acting based on knowledge at hand

Specifying the task environment

PEAS (Performance, Environment, Actuators, Sensors)

Agent Type	Performance Measure	Environment	Actuators	Sensors
Taxi driver	Safe, fast, legal, comfortable trip, maximize profits	Roads, other traffic, pedestrians, customers	Steering, accelerator, brake, signal, horn, display	Cameras, sonar, speedometer, GPS, odometer, accelerometer, engine sensors, keyboard

Examples

Agent Type	Performance Measure	Environment	Actuators	Sensors
Medical diagnosis system	Healthy patient, reduced costs	Patient, hospital, staff	Display of questions, tests, diagnoses, treatments, referrals	Keyboard entry of symptoms, findings, patient's answers
Satellite image analysis system	Correct image categorization	Downlink from orbiting satellite	Display of scene categorization	Color pixel arrays
Part-picking robot	Percentage of parts in correct bins	Conveyor belt with parts; bins	Jointed arm and hand	Camera, joint angle sensors
Refinery controller	Purity, yield, safety	Refinery, operators	Valves, pumps, heaters, displays	Temperature, pressure, chemical sensors
Interactive English tutor	Student's score on test	Set of students, testing agency	Display of exercises, suggestions, corrections	Keyboard entry

Environments

Environment examples

Task Environment	Observable	Agents	Deterministic	Episodic	Static	Discrete
Crossword puzzle Chess with a clock	Fully Fully	Single Multi	Deterministic Deterministic	Sequential Sequential	Static Semi	Discrete Discrete
Poker Backgammon	Partially Fully	Multi Multi	Stochastic Stochastic	Sequential Sequential	Static Static	Discrete Discrete
Taxi driving Medical diagnosis	Partially Partially	Multi Single	Stochastic Stochastic	Sequential Sequential	Dynamic Dynamic	Continuous Continuous
Image analysis Part-picking robot	Fully Partially	Single Single	Deterministic Stochastic	Episodic Episodic	Semi Dynamic	Continuous Continuous
Refinery controller Interactive English tutor	Partially Partially	Single Multi	Stochastic Stochastic	Sequential Sequential	Dynamic Dynamic	Continuous Discrete

Episodic: The next episode does not depend on the actions taken in previous episodes

Sequential: The current decision could affect all future decisions

Deterministic: The next state of the environment is completely determined by the current state and the action

Stochastic: Actions are characterized by their possible outcomes

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Accessible??</u>	Yes	Yes	No	No
<u>Deterministic??</u>	Yes	No	Partly	No
<u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>	Yes	Yes	Yes	No

The environment type largely determines the agent design

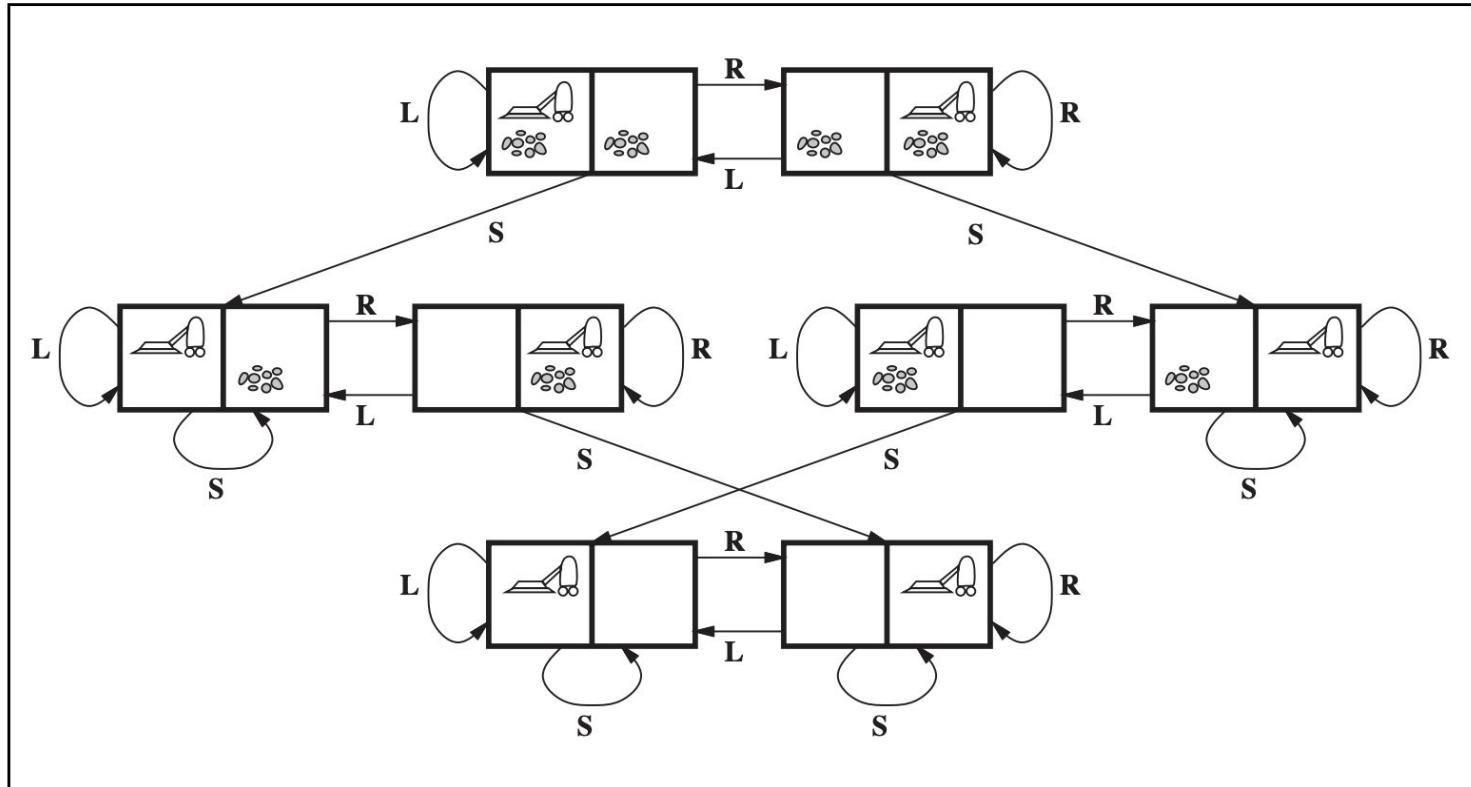
The real world is (of course) inaccessible, stochastic, sequential, dynamic, continuous

Solving problems by Search

Problem solving agents

- Goal formulation
- Problem formulation
- Environment
- Actions

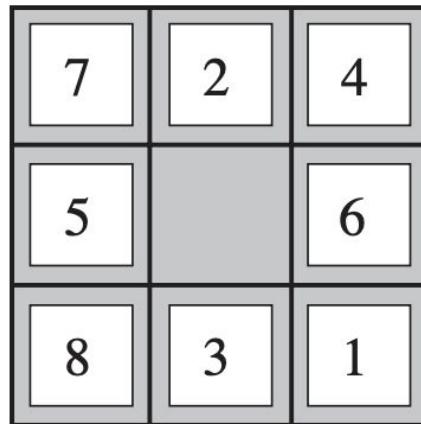
Example - *Vacuum cleaner*



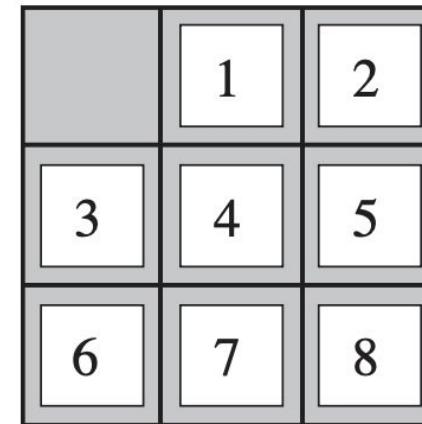
Vacuum cleaner - *Problem formulation*

- **States:** The state is determined by both the agent location and the dirt locations. The agent is in one of two locations, each of which might or might not contain dirt. Thus, there are $2 \times 2^2 = 8$ possible world states. A larger environment with n locations has $n \cdot 2^n$ states.
- **Initial state:** Any state can be designated as the initial state.
- **Actions:** In this simple environment, each state has just three actions: *Left*, *Right*, and *Suck*. Larger environments might also include *Up* and *Down*.
- **Transition model:** The actions have their expected effects, except that moving *Left* in the leftmost square, moving *Right* in the rightmost square, and *Sucking* in a clean square have no effect. The complete state space is shown in Figure 3.3.
- **Goal test:** This checks whether all the squares are clean.
- **Path cost:** Each step costs 1, so the path cost is the number of steps in the path.

Example - *8-puzzle*



Start State

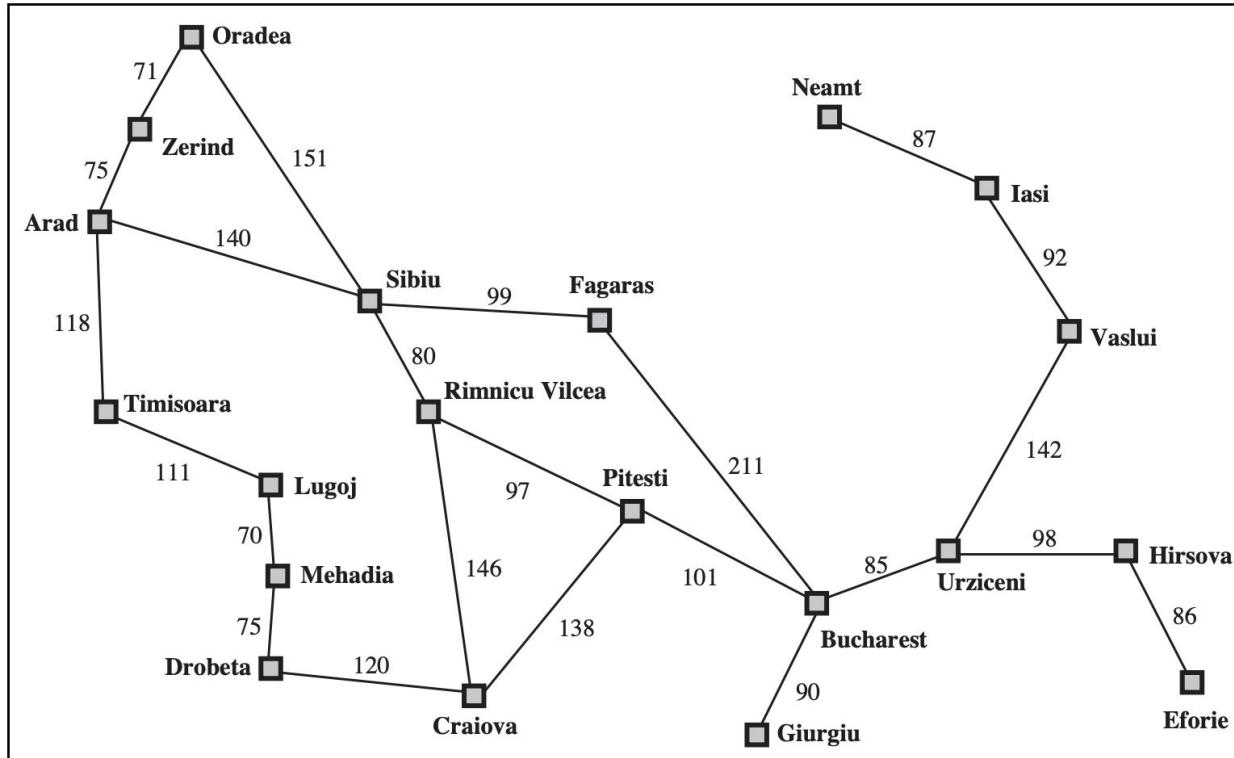


Goal State

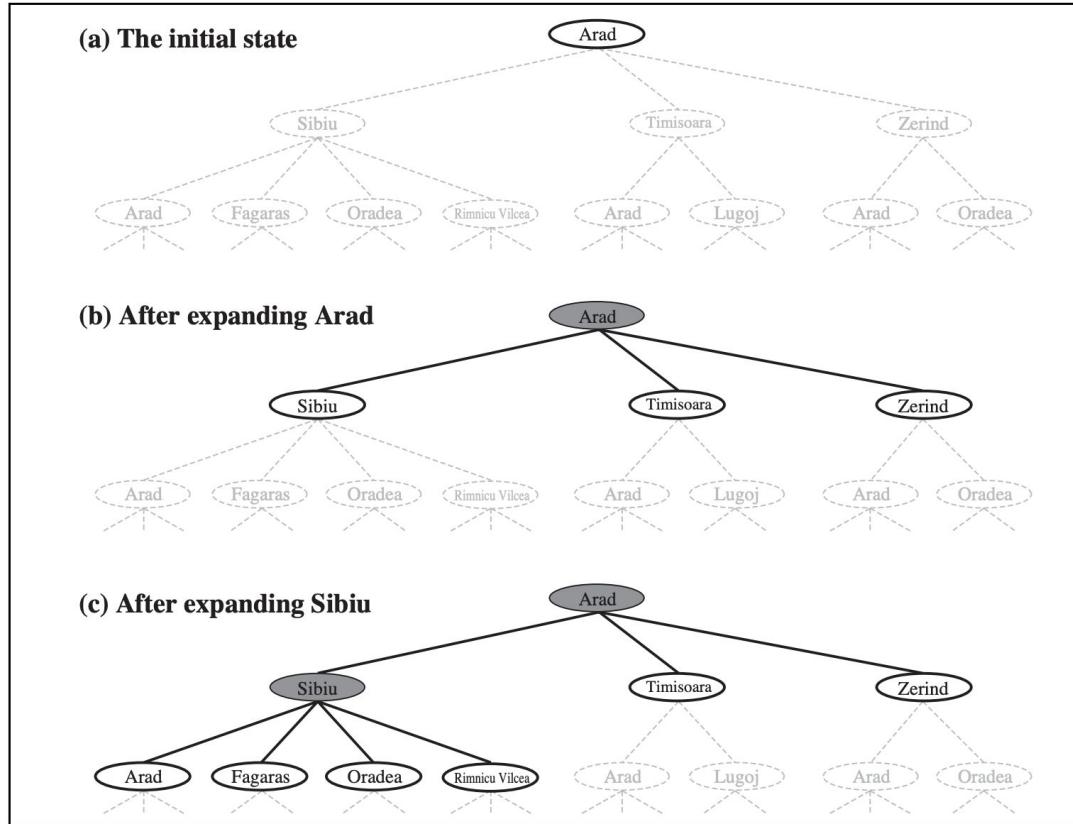
8-puzzle

- **States:** A state description specifies the location of each of the eight tiles and the blank in one of the nine squares.
- **Initial state:** Any state can be designated as the initial state. Note that any given goal can be reached from exactly half of the possible initial states
- **Actions:** The simplest formulation defines the actions as movements of the blank space *Left*, *Right*, *Up*, or *Down*. Different subsets of these are possible depending on where the blank is.
- **Transition model:** Given a state and action, this returns the resulting state; for example, if we apply *Left* to the start state in Figure 3.4, the resulting state has the 5 and the blank switched.
- **Goal test:** This checks whether the state matches the goal configuration shown in Figure 3.4. (Other goal configurations are possible.)
- **Path cost:** Each step costs 1, so the path cost is the number of steps in the path.

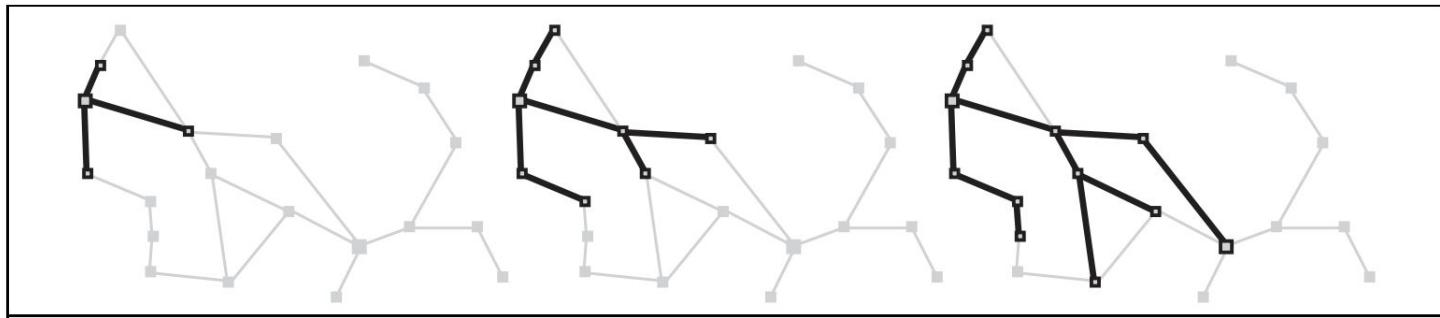
Example - Roadtrip from Arad to Bucharest



Roadtrip

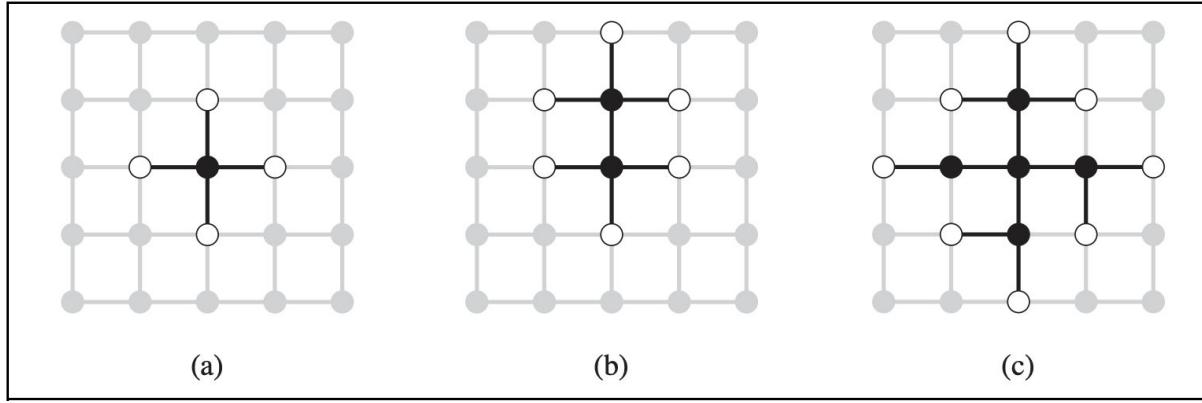


Roadtrip - Tree Search



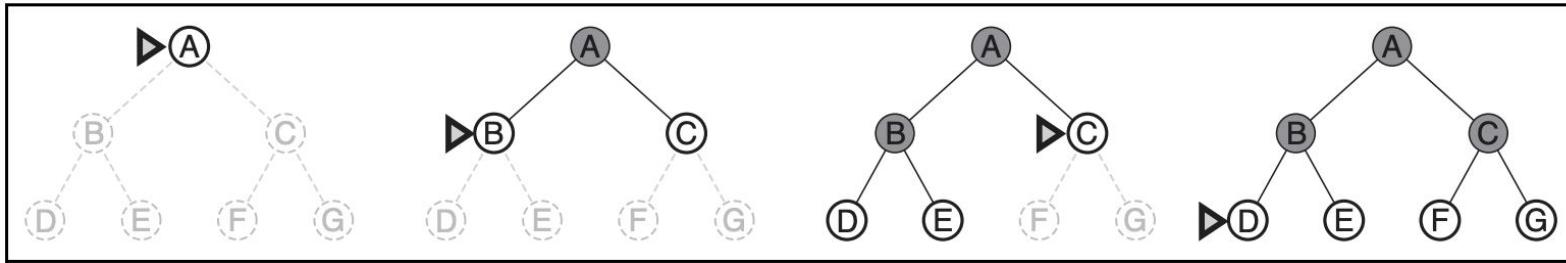
```
function TREE-SEARCH(problem) returns a solution, or failure
    initialize the frontier using the initial state of problem
    loop do
        if the frontier is empty then return failure
        choose a leaf node and remove it from the frontier
        if the node contains a goal state then return the corresponding solution
        expand the chosen node, adding the resulting nodes to the frontier
```

Roadtrip - Graph Search



```
function GRAPH-SEARCH(problem) returns a solution, or failure
    initialize the frontier using the initial state of problem
    initialize the explored set to be empty
    loop do
        if the frontier is empty then return failure
        choose a leaf node and remove it from the frontier
        if the node contains a goal state then return the corresponding solution
        add the node to the explored set
        expand the chosen node, adding the resulting nodes to the frontier
        only if not in the frontier or explored set
```

Uninformed Search - Breadth First Search

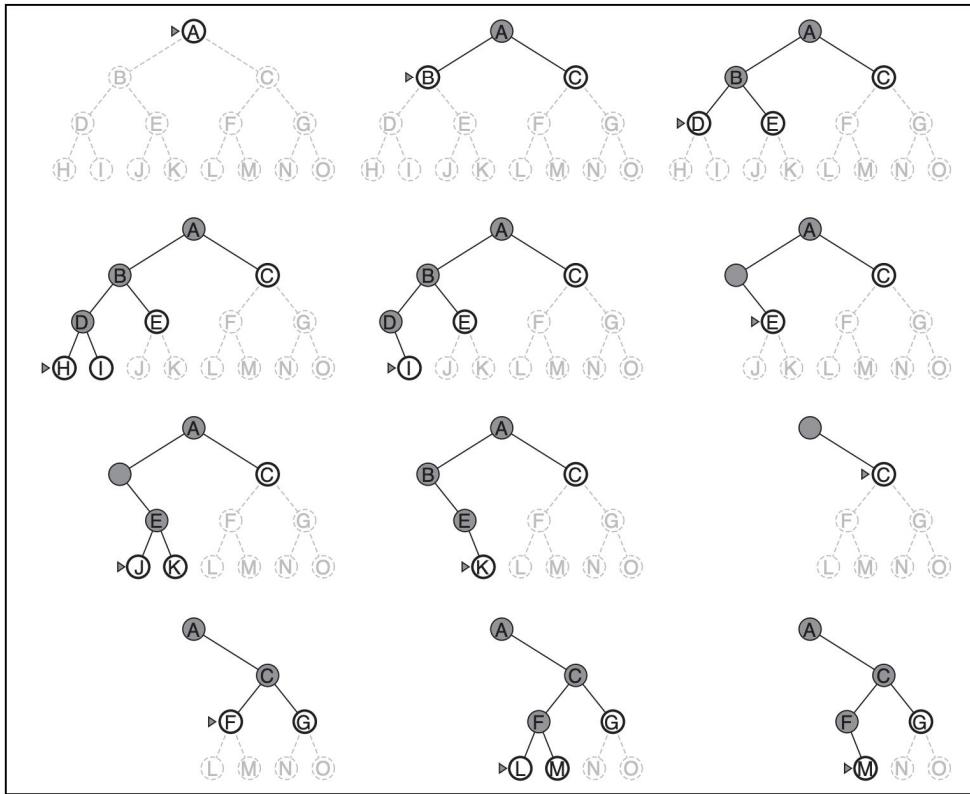


```
function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier  $\leftarrow$  a FIFO queue with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier  $\leftarrow$  INSERT(child, frontier)
```

Time and space complexity of BFS

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

Uninformed Search - Depth first search



Informed Search (Heuristic search)

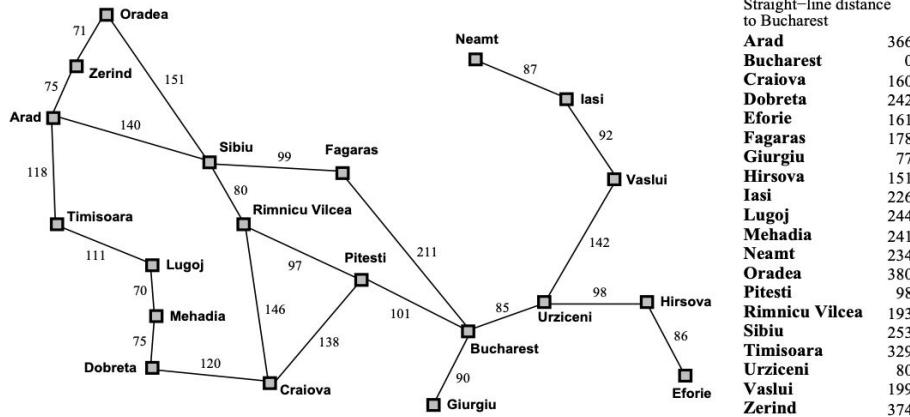
Heuristic Search

- **Heuristic or informed search** exploits additional knowledge about the problem that helps direct search to more promising paths.
- A **heuristic function**, $h(n)$, provides an estimate of the cost of the path from a given node to the closest goal state.
Must be zero if node represents a goal state.
 - Example: Straight-line distance from current location to the goal location in a road navigation problem.
- Many search problems are NP-complete so in the worst case still have exponential time complexity; however a good heuristic can:
 - Find a solution for an average problem efficiently.
 - Find a reasonably good but not optimal solution efficiently.

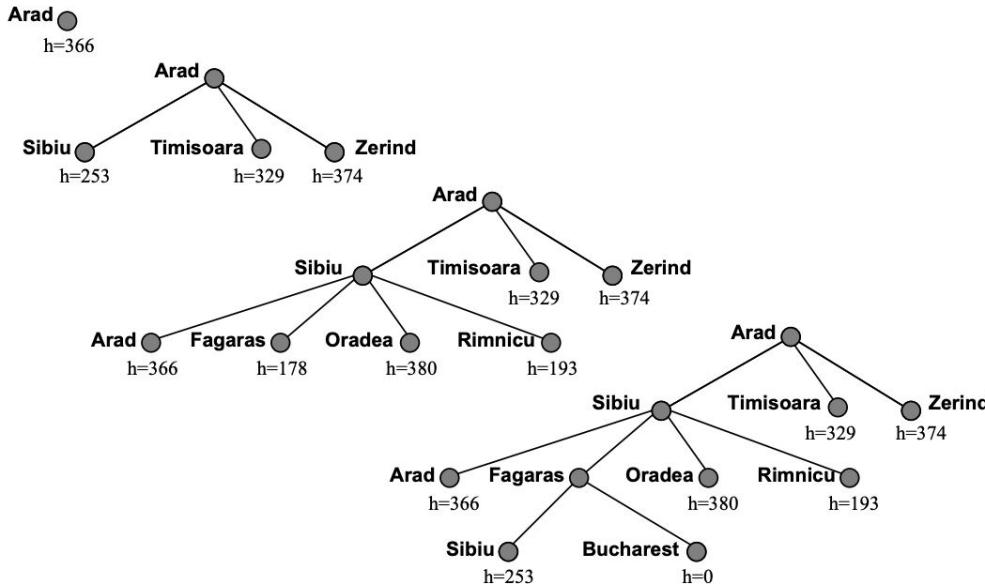
Best-First Search

- At each step, best-first search sorts the queue according to a heuristic function.

function BEST-FIRST-SEARCH(*problem*, EVAL-FN) **returns** a solution sequence
inputs: *problem*, a problem
Eval-Fn, an evaluation function
Queueing-Fn \leftarrow a function that orders nodes by EVAL-FN
return GENERAL-SEARCH(*problem*, *Queueing-Fn*)



Best-First Example



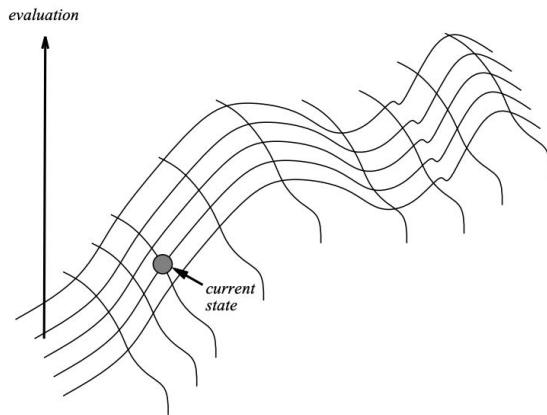
- Does not find shortest path to goal (through Rimnicu) since it is only focused on the cost remaining rather than the total cost.

Best-First Properties

- Not complete, may follow infinite path if heuristic rates each state on such a path as the best option. Most reasonable heuristics will not cause this problem however.
- Worst case time complexity is still $O(b^m)$ where m is the maximum depth.
- Since must maintain a queue of all unexpanded states, space-complexity is also $O(b^m)$.
- However, a good heuristic will avoid this worst-case behavior for most problems.

Hill-Climbing

- Beam search with a beam-width of 1 is called **hill-climbing**.
- Pursues locally best option at each point, i.e. the best successor to the current best node.
- Subject to local maxima, plateaux, and ridges.



Minimizing Total Path Cost: A* Search

- A* combines features of uniform cost search (complete, optimal, inefficient) with best-first (incomplete, non-optimal, efficient).
- Sort queue by estimated total cost of the completion of a path.

$$f(n) = g(n) + h(n)$$

- If the heuristic function always underestimates the distance to the goal, it is said to be **admissible**.
- If h is admissible, then $f(n)$ never overestimates the actual cost of the best solution through n .

Other possible search options

- Genetic algorithm
- Djikstra's algorithm (for route planning)
- Simulated annealing
- Hill Climbing